# Regularized AFT models

Anuj Khare

This has been adapted from the Survival package vignette.

## Model definition

**Inputs:** $X_{n*p}$, $T_{n*2}$, density $f$...
**Parameters:** Coefficients $\beta = [\beta_0, \beta_1..., \beta_p]^T$, Scale $\sigma$
where, $\beta_0$ is the intercept.

The type of censoring is determined as follows:

$$\begin{cases} \text{Left censoring} & \text{if: } -\infty = \underline{t}_i \text{ , } \bar{t}_i < \infty \\ \text{Right censoring} & \text{if: } -\infty < \underline{t}_i \text{ , } \bar{t}_i = \infty \\ \text{Interval censoring} & \text{if: } -\infty < \underline{t}_i \neq \bar{t}_i < \infty \\ \text{No censoring} & \text{if: } -\infty < \underline{t}_i = \bar{t}_i < \infty \end{cases} \tag{1}$$

Define the transformed output,

$$y_i = trans(T_i) \tag{2}$$

where $trans$ depends on the distribution, and is log for the log-gaussian, log-logistic distributions, and so on.

The model is defined by the equation:

$$y = X\beta + \sigma\epsilon \tag{3}$$

where, $\epsilon \sim f$. Thus,

$$e_i = \frac{y_i - x_i^T\beta}{\sigma} \sim f \tag{4}$$

We define the elastic net (L1 + L2) penalty as follows:

$$\lambda P_\alpha(\beta) = \lambda(\alpha\|\beta\|_1 + 1/2(1-\alpha)\|\beta\|_2^2) \tag{5}$$

Our objective is to maximize the penalized, scaled log likelihood:

$$\hat{\beta} = argmax_\beta \left(\frac{1}{n} l(\beta) - \lambda P_\alpha(\beta)\right) \tag{6}$$

## Likelihood

For calculating likelihood, in the observations with no censoring, the pdf is used, and in censored observations, the cdf is used. Hence, the likelihood is given as:

$$lik = \left(\prod_{exact} f(e_i)/\sigma\right)\left(\prod_{right} 1 - F(e_i)\right)\left(\prod_{left} F(e_i)\right)\left(\prod_{interval} F(e_i^u) - F(e_i^l)\right) \tag{7}$$

"Exact", "left", "right", and "interval" refer to uncensored, left censored, right censored and interval censored observations respectively, and $F$ is the cdf of the distribution. $e_i^u$, and $e_i^l$ are upper and lower endpoints for interval censored data.

Hence the log likelihood is given as:

$$l(\beta) = \sum_{exact} g_1(e_i) - \log(\sigma) + \sum_{right} g_2(e_i) + \sum_{left} g_3(e_i) + \sum_{interval} g_4(e_i^l, e_i^u) \tag{8}$$

$g_1 = \log(f)$, $g_2 = \log(1 - F)$, $g_3 = \log(F)$, $g_4(e_i^l, e_i^u) = \log(F(e_i^u) - F(e_i^l))$.

## Score and Hessian

Derivatives of the LL with respect to the regression parameters are:

$$\frac{\partial l(\beta)}{\partial \beta_j} = \sum_{i=1}^{n} \frac{\partial g}{\partial \eta_i}\frac{\partial \eta_i}{\partial \beta_j} = \sum_{i=1}^{n} x_{ij}\frac{\partial g}{\partial \eta_i} \tag{9}$$

$$\frac{\partial^2 l(\beta)}{\partial \beta_j \beta_k} = \sum_{i=1}^{n} x_{ij}x_{ik}\frac{\partial^2 g}{\partial \eta_i^2} \tag{10}$$

where $\eta_i = x_i^T\beta$ is the vector of linear predictors.

Define $\mu_i = \frac{\partial g}{\partial \eta_i}$, where $g$ is one of $g_1$ to $g_4$ depending on type of censoring in the $i^{th}$ observation, and $\mu = [\mu_1, ...\mu_n]^T$. Then, partial derivative of log-likelihood is given as:

$$\frac{\partial l(\beta)}{\partial \beta_j} = \sum_{i=1}^{n} x_{ij}\mu_i \tag{11}$$

Hence, the score (gradient of log likelihood) is given as:

$$S = \nabla_\beta l(\beta) = X^T \mu = \sum_{i=1}^{n} \mu_i \overline{x}_i \tag{12}$$

The hessian can be written as:

$$H = \sum_{i=1}^{n} \overline{x}_i \overline{x}_i^T \frac{\partial^2 g}{\partial \eta_i^2} = \sum_{i=1}^{n} \overline{x}_i \overline{x}_i^T w_i \tag{13}$$

Define $W = diag(w_1, ...w_n)$.

$$H = X^T W X \tag{14}$$

## Taylor approximation for log-likeihood

A 2-step Taylor series centered at $\widetilde{\beta}$ is given as:
where, define the working response $\widetilde{z} = X\widetilde{\beta} + \widetilde{W}^{-1}\widetilde{\mu}$. Here, the tilde denotes that the respective values are evaluated using the parameters from the previous step.
Hence, the log likelihood can be approximated centered at $\widetilde{\beta}$ as:

$$l(\beta) =<> \tag{15}$$

This algorithm is the iteratively reweighted least squares (IRLS), since at each iteration we solve a weighted least squares problem.

## Coordinate descent

Hence, at each step we are solving a penalized weighted least squares problem, which is a minimizer of (using the scaled approximate log-likelihood):

$$M = \frac{1}{2n} \sum_{i=1}^{n} \widetilde{w}_i (\widetilde{z}_i - \overline{x}_i^T \beta)^2 + \lambda P_\alpha^(\beta) \tag{16}$$

The subderivative of the optimization objective is given as:

$$\frac{\partial M}{\partial \beta_k} = \frac{1}{n} \sum_{i=1}^{n} \widetilde{w}_i x_{ik} (\widetilde{z}_i - \overline{x}_i^T \beta) + \lambda\alpha \ \text{sgn}(\beta_k) + \lambda(1-\alpha)\beta_k \tag{17}$$

where, sgn($\beta_k$) is 1 if $\beta_k > 1$, -1 if $\beta_k < 0$ and 0 if $\beta_k = 0$.
Using the subderivative, three cases of solutions for $\beta_k$ may be obtained.
The solution is given by:

$$\hat{\beta}_k = \frac{S\left(\frac{1}{n}\sum_{i=1}^{n} \widetilde{w}_i x_{ik}\left[\widetilde{z}_i - \sum_{j\neq k} x_{ij}\beta_j\right], \lambda\alpha\right)}{\frac{1}{n}\sum_{i=1}^{p} \widetilde{w}_i x_{ik}^2 + \lambda(1-\alpha)} \tag{18}$$

where, S is the soft thresholding operator, and $w_i$ and $z_i$ are given in 13 and 16 respectively.
The intercept is not regularized, and hence can be calculated as:

$$\hat{\beta}_0 = \frac{\frac{1}{n}\sum_{i=1}^{n} \widetilde{w}_i\left[\widetilde{z}_i - \sum_{j\neq 0} x_{ij}\beta_j\right]}{\frac{1}{n}\sum_{i=1}^{p} \widetilde{w}_i} \tag{19}$$

The coordinate descent algorithm works by cycling through each $\beta_j$ in turn, keeping the others constant, and using the above estimate to calculate the optimal value $\hat{\beta}_j$.

After each update cycle for $\beta$, the scale parameter $\sigma$ is updated once using a Newton step:

$$\sigma_{new} = \sigma_{old} - \left(\frac{\partial l^2(\sigma)}{\partial \sigma^2}\right)^{-1}\left(\frac{\partial l(\sigma)}{\partial \sigma}\right) \tag{20}$$

This is repeated until convergence of both $\beta$ and $\sigma$. Note that we have ignored the off-diagonal entries in the Hessian for the scale parameter.

## Pathwise solution

This section is borrowed from section 2.3 of [3]. The iregnet function will return solutions for an entire path of vaules of $\lambda$, for a fixed $\alpha$. We begin with $\lambda$ sufficiently large to set the solution $\beta = 0$, and decrease $\lambda$ until we arrive near the unregularized solution. The solutions for each value of $\lambda$ are used as the initial estimates of $\beta$ for the next $\lambda$ value. This is known as warm starting, and makes the algorithm efficient and stable. To choose initial value of $\lambda$, we use Equation 18, and notice that for $\frac{1}{n}\sum_{i=1}^{n} w_i(0)x_{ij}z(0)_i < \alpha\lambda$ for all $j$, then $\beta = 0$ minimizes the objective 6. Thus,

$$\lambda_{max} = max_j \frac{1}{n\alpha}\sum_{i=1}^{n} w_i(0)x_{ij}z(0)_i \tag{21}$$

We will set $\lambda_{min} = \epsilon\lambda_{max}$ , and compute solutions over a grid of $m$ values, where $\lambda_j = \lambda_{max}(\lambda_{min}/\lambda_{max})^{j/m}$ for $j = 0, .., m$.

## Algorithm

The algorithm to be followed for fitting the distribution is:

Transform output variable $y$ using log transformation ;
Calculate $\lambda_{max}$ using equation 21, and set $\widetilde{\beta} = 0$, $\widetilde{\eta} = 0$ ;
Calculate $\lambda_{min}$ and a grid of $m$ $\lambda$ values ;
**foreach** $\lambda_j$ *in* $j = m, ..., 0$ **do**
    **repeat**
        Compute $\widetilde{w}_i$ and $\widetilde{z}_i$ ;
        Find $\hat{\beta}$ by solving the penalized weighted least square problem defined in equation 6 using coordinate descent ;
        Set $\widetilde{\beta} = \hat{\beta}$ ;
    **until** *convergence of* $\hat{\beta}$;
    Set $\widetilde{\beta} = \hat{\beta}$ , $\widetilde{\eta} = X\widetilde{\beta}$ ;
**end**

**Algorithm 1:** Overall optimization algorithm

## Scale parameter

So far, I have ignored the $\sigma$ parameter from the calculations and equations. This is only reasonable if we treat $\sigma$ as fixed. However, in other cases, $\sigma$ needs to estimated along with the parameters $\beta$, by using the derivatives as listed below.

## Derivatives

Iterations are done with respect to $\log(\sigma)$ to prevent numerical underflow.

$$
\begin{aligned}
\frac{\partial g_1}{\partial \eta} &= -\frac{1}{\sigma}\left[\frac{f'(z)}{f(z)}\right] \\
\frac{\partial g_4}{\partial \eta} &= -\frac{1}{\sigma}\left[\frac{f(z^u)-f(z^l)}{F(z^u)-F(z^l)}\right] \\
\frac{\partial^2 g_1}{\partial \eta^2} &= -\frac{1}{\sigma^2}\left[\frac{f''(z)}{f(z)}\right] - (\partial g_1/\partial \eta) \\
\frac{\partial^2 g_4}{\partial \eta^2} &= -\frac{1}{\sigma^2}\left[\frac{f'(z^u)-f'(z^l)}{F(z^u)-F(z^l)}\right] - (\partial g_4/\partial \eta)^2 \\
\frac{\partial g_1}{\partial \log \sigma} &= -\left[\frac{z f'(z)}{f(z)}\right] \\
\frac{\partial g_4}{\partial \log \sigma} &= -\left[\frac{z^u f(z^u)-z^l f(z^l)}{F(z^u)-F(z^l)}\right] \\
\frac{\partial^2 g_1}{\partial (\log \sigma)^2} &= \left[\frac{z^2 f''(z)+z f'(z)}{f(z)}\right] - (\partial g_1/\partial \log \sigma)^2 \\
\frac{\partial^2 g_4}{\partial (\log \sigma)^2} &= \left[\frac{(z^u)^2 f'(z^u)-(z^l)^2 f'(z^l)}{F(z^u)-F(z^l)}\right] - (\partial g_1/\partial \log \sigma)(1+\partial g_1/\partial \log \sigma) \\
\frac{\partial^2 g_1}{\partial \eta \partial \log \sigma} &= \left[\frac{z f''(z)}{\sigma f(z)}\right] - (\partial g_1/\partial \eta)(1+\partial g_1/\partial \log \sigma) \\
\frac{\partial^2 g_4}{\partial \eta \partial \log \sigma} &= \left[\frac{z^u f'(z^u)-z^l f'(z^l)}{\sigma[F(z^u)-F(z^l)]}\right] - (\partial g_4/\partial \eta)(1+\partial g_4/\partial \log \sigma)
\end{aligned}
\tag{22}
$$

Derivatives for $g_2$ can be obtained by setting $z_u$ to inf in the equations for $g_4$, and similarly for $g_3$.
The distribution specific values of $f(z)$, etc. are omitted.

## Subgradient of Cost

The cost to be minimized is the negative of the penalized, scaled log-likelihood:

$$J(\beta) = \left( -\frac{1}{n} l(\beta) + \lambda P_\alpha(\beta) \right) \tag{23}$$

$$\hat{\beta} = argmin_\beta \left( -\frac{1}{n} l(\beta) + \lambda P_\alpha(\beta) \right) \tag{24}$$

The subderivative of the cost is given as:

$$\nabla_\beta J = -\frac{1}{n} S(\beta) + \lambda \alpha \text{sgn}(\beta) + \lambda(1-\alpha)\beta \tag{25}$$

where, sgn($\beta$) is calculated element-wise on the vector. S is the score as given in 12.

The closeness of the degree 1, 2, and inf norms of the subderivate to zero can be used as a mteric for judging the optimality of the obtained solutions.

# Bibliography

[1] Survival - Terry M Therneau

[2] Machine Learning: A Probabilistic Perspective - Kevin Murphy

[3] Regularization Paths for Cox's Proportional Hazards Model via Coordinate Descent - Simon, Friedman, Hastie, Tibshirani

[4] AFT - TD Hocking